



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



# Writing GPU Kernels

Ben Cumming, CSCS  
February 25, 2017



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

# Going Parallel : Kernels and Threads

---

# Threads and Kernels

- **Threads** are streams of execution, run simultaneously on GPU (1000s)
- A **kernel** is the function run by each thread
- CUDA provides language support for:
  - writing kernels;
  - launching many threads to execute a kernel in parallel.
- CUDA hides the low-level details of launching threads.

## The process for developing CUDA kernels

1. Formulate algorithm in terms of parallel work items.
2. Write a kernel implementing a work item on one thread.
3. Launch the kernel with the required number of threads.

# Scaled Vector Addition (axpy)

We have used CUBLAS to perform scaled vector addition:

$$\mathbf{y} = \mathbf{y} + \alpha \mathbf{x}$$

- $\mathbf{x}$  and  $\mathbf{y}$  are vectors of length  $n$ ;

$$x, y \in \mathbb{R}^n$$

- $\alpha$  is scalar.

$$\alpha \in \mathbb{R}$$

Applying **axpy** requires  $n$  operations:

$$y_i \leftarrow y_i + a * x_i, \quad i = 0, 1, \dots, n - 1$$

which can be performed **independently** and **in any order**.

axpy implemented on CPU with a loop

```
void axpy(double *y, const double *x, double a, int n) {  
    for(int i=0; i<n; ++i)  
        y[i] = y[i] + a*x[i];  
}
```

# Kernels

A **kernel** defines the work item for a single thread

- The work is performed by many threads executing the same kernel **simultaneously**.
- Conceptually corresponds to the inner part of a loop for BLAS1 operations like **axpy**.

host : add two vectors

```
void add_cpu(int *a, int *b, int n){  
    for(auto i=0; i<n; ++i)  
        a[i] = a[i] + b[i];  
}
```

CUDA : add two vectors

```
--global--  
void add_gpu(int *a, int *b, int n){  
    auto i = threadIdx.x;  
    a[i] = a[i] + b[i];  
}
```

- `--global--` keyword indicates a kernel
- `threadIdx` used to find unique id of each thread

# Launching a kernel

- Host code launches a kernel on the GPU **asynchronously**.
- CUDA provides the “tripple chevron” `<<<_,_>>>` syntax for launching a kernel.

host : add two vectors

```
auto n = 1024;  
auto a = host_malloc<int>(n);  
auto b = host_malloc<int>(n);  
add_cpu(a, b, n);
```

CUDA : add two vectors

```
auto n = 1024;  
auto a = device_malloc<int>(n);  
auto b = device_malloc<int>(n);  
add_gpu<<<1,n>>>(a, b, n);
```

- `add_gpu<<<1, num_threads>>>(args... )` launches the kernel  
`add_gpu` with `num_threads` parallel threads.

# Exercise: My First Kernel

Open `axpy/axpy.cu`

1. Write a kernel that implements `axpy` for `double`
  - `axpy_kernel(double *y, double *x, double a, int n)`
  - **extra:** can you write a C++ templated version for any type?
2. launch the kernel (look for `TODO`)
3. Compile the test and run
  - it will pass with no errors on success
  - first try with small vectors of size 8
  - try increasing launch size... what happens?
4. **extra:** can you extend the kernel to work for larger arrays?



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

# Scaling Up : Thread Blocks

---



In the **axpy** exercises we were limited to 1024 threads for a kernel launch

- but we need to scale beyond 1024 threads for the **massive parallelism** we were promised!

## Thread blocks and grids

kernels are executed in groups of threads called **thread blocks**



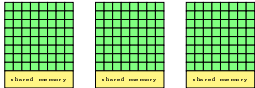
- the launch configuration `axpy<<grid_dim, block_dim>>>(...)`
  - launch a **grid** of `grid_dim` **blocks**
  - each **block** has `block_dim` **threads**
  - for a total of `grid_dim × block_dim` threads
- previously we launched just one thread block

```
axpy<<1, n>>>(...)
```

# Why the additional complexity?

## Coordination between threads doesn't scale:

- Threads in a block can synchronize and share resources
- This does not scale past a certain number of cores/threads
- EACH P100 GPU **streaming multiprocessor** (SMX) has 64 CUDA cores, and can run 2028 threads
- Threads in a block run on the same SMX, with shared resources and thread cooperation
- Work is broken into blocks, which are distributed over the 56 SMXs on the GPU.

concept	hardware	
thread		<ul style="list-style-type: none"> <li>■ each thread executed on one core</li> </ul>
block		<ul style="list-style-type: none"> <li>■ block executed on 1 SMX</li> <li>■ multiple blocks per SMX if sufficient resources</li> <li>■ threads in a block share SMX resources</li> </ul>
grid		<ul style="list-style-type: none"> <li>■ kernel is executed in grid of blocks</li> <li>■ blocks distributed over SMXs</li> <li>■ multiple kernels can run at same time</li> </ul>

# Calculating thread indexes

A kernel has to calculate the index of its work item

- In `axpy` we used `threadIdx.x` for the index.
- With multiple blocks, we need more information, which is available in the following **magic variables**:

<code>gridDim</code>	: total number of blocks in the grid
<code>blockDim</code>	: number of threads in a thread block
<code>blockIdx</code>	: index of block <code>[0, gridDim-1]</code>
<code>threadIdx</code>	: index of thread in thread block <code>[0, blockDim-1]</code>


# Calculating thread indexes

Consider accessing an array of length 24 with 8 threads per block. The **dimensions** of the kernel launch are:

- `blockDim.x == 8` (8 threads/block)
- `gridDim.x == 3` (3 blocks)

We calculate the index for our thread using the formula

```
auto index = threadIdx.x + blockDim.x*blockIdx.x
```



# Calculating grid dimensions

The number of thread blocks and the number of threads per block are parameters for the kernel launch:

```
kernel<<<blocks, threads_per_block>>>(...)
```

Remember to guard against overflow when the number of work items is not divisible by the thread block size

## vector addition with multiple blocks

```
__global__  
void add_gpu(int *a, int *b, int n){  
    auto i = threadIdx.x + blockIdx.x*blockDim.x;  
    if(i<n) { // guard against access off end of arrays  
        a[i] += b[i];  
    }  
}  
  
// in main()  
auto block_size = 512;  
auto num_blocks = (n + (block_size-1)) / block_size;  
add_gpu<<<num_blocks, block_size>>>(a, b, n);
```

# Calculating grid dimensions

We have to take care when calculating the number of blocks in the grid, i.e. `blocks`:

```
kernel<<<blocks, threads_per_block>>>(...)
```

Most likely, the number of work items `n` is not a multiple of `threads_per_block`

- some threads in the last thread block will be idle.

## Calculating grid dimensions

```
// in main()
auto block_size = 512;
auto num_blocks = (n + (block_size-1)) / block_size;
add_gpu<<<num_blocks, block_size>>>(a, b, n);
```

# How many threads per block?

The number of threads per block has an impact on performance

- The optimal number depends on resources required by the kernel (registers, shared memory, computational intensity, etc).

The short answer is 64 or 128 on P100.

- Wait for the P100 deep dive session for more information...




# Exercise: Blocks

Open `cuda/exercises/axpy/axpy.cu` from the last exercise

1. Extend the `axpy` kernel for arbitrarily large input arrays (any `n`)
2. Update the call site to calculate the grid configuration
3. Compile the test and run
  - it will pass with no errors on success
4. Experiment with varying the size of the arrays (scaling)
  - start small and increase
5. finish the `newton.cu` example
  - how do the h2d, d2h and kernel timings compare?
6. **extra:** Compare scaling with the `axpy_omp` benchmark
7. **extra:** Experiment with varying the block size

## Exercise: Results



The GPU is a throughput device:

- CUDA breaks even for  $n \geq 2^{14} \approx 16,000$
- requires  $2^{21} \approx 2,000,000$  to gain “full”  $7\times$  speedup

You have to provide enough parallelism to exploit many cores